

Limitations of Anonymous, Self-Reported Admissions Data

Joshua Latz (jlatz1) · EN 605.256 Modern Software Concepts in Python · Module 3

What the database contains is shaped as much by what applicants left blank as by what they entered. Every entry has a status, term, and degree, but only 60.4% report a GPA and just 8.0% report any GRE score, and the gaps are not random. Among Fall 2026 entries, those with a GPA report an acceptance 40.71% of the time, against 31.73% for those without one, so averages over reported values describe a group tilted toward acceptances. The values that are reported are also unreliable. The field the site labels GRE holds combined scores for 1,406 entries and quantitative scores for only 954, alongside impossible values such as 999, so the average GRE Quantitative score of 165.83 rests on 3.2% of the database and on validity ranges I chose. Range checks catch only impossible values: a plausible typo, or a Verbal score entered as Quantitative, passes unnoticed. The database mathematically contains an average of 165.83 over 954 entries. It does not tell us the average score of the people who reported a GRE, let alone of everyone who applied.

The deeper limitation is who chooses to post. Nothing identifies a person, so the database counts posts, not applicants: someone applying to ten programs may appear ten times, and Question 1's 29,576 Fall 2026 entries is not a number of applicants. The posts also show clear self-reporting bias. Fall 2025 entries, posted in 2026 after that cycle had largely finished, report acceptances at 47.92%, more than ten points above Fall 2026's 37.10%, consistent with late posters being those with good news. Reported GRE averages run about one standard deviation above the means ETS publishes for all test takers, 165.83 against 157.62 for Quantitative (ETS, GRE General Test Interpretive Data, Table 1A). Part of that gap is real, since applicants to competitive programs are not typical test takers, but the data cannot separate it from strong scorers being likelier to report. So the database mathematically contains a 37.33% share of posts marked Accepted and nearly equal PhD acceptance shares for American and International entries. It does not show that 37% of applicants are admitted, or that nationality has no bearing on admission. It supports only descriptions of what self-selected posters reported, and evidence, like the results above, of how that reporting is biased.